

Some ways to get data into Databricks

- manually upload files to volume
- create incremental ingestion with Lakeflow Connect
 - fully-managed connectors
 - standard connectors

Some ways to get data into Databricks

- manually upload files to volume
- create incremental ingestion with Lakeflow Connect
 - fully-managed connectors
 - standard connectors

only raw files
inside Databricks,
you create subsequent
objects

Some ways to get data into Databricks

- manually upload files to volume
- create incremental ingestion with Lakeflow

Connect

- fully-managed connectors

SaaS applications:
Meta Ads,
Salesforce,
ServiceNow etc

Databases:
MySQL,
MS SQL Server,
PostgreSQL

use simple UI to create
connection
+ ingestion pipeline
+ destination

- standard connectors

Some ways to get data into Databricks

- manually upload files to volume
- create incremental ingestion with Lakeflow Connect

- fully-managed connectors

SaaS applications:
Meta Ads,
Salesforce,
ServiceNow etc

Databases:
MySQL,
MS SQL Server,
PostgreSQL

- standard connectors

Other data sources:
SFTP Servers,
Apache Kafka etc

code-based
with more customization

Common table & view types

- managed table
- streaming table
- view
- materialized view

Common table & view types

- managed table: Databricks fully manages both the data storage and metadata
- streaming table
- view
- materialized view

Common table & view types

- managed table: Databricks fully manages both the data storage and metadata
- streaming table: a table that is only updated with new data
- view
- materialized view

Common table & view types

- managed table: Databricks fully manages both the data storage and metadata
- streaming table: a table that is only updated with new data
- view: a virtual table defined by a query that is computed each time it is accessed
- materialized view

Common table & view types

- managed table: Databricks fully manages both the data storage and metadata
- streaming table: a table that is only updated with new data
- view: a virtual table defined by a query that is computed each time it is accessed
- materialized view: a view where the query result is physically stored

ETL pipeline in Databricks

- can be built with Lakeflow Spark Declarative Pipelines (SDP) or Apache Spark

ETL pipeline in Databricks

- can be built with Lakeflow Spark Declarative Pipelines (SDP) or Apache Spark
- simplified steps with Lakeflow SDP:

